

```
In [2]: #1. Download the employee attrition dataset
#https://www.kaggle.com/datasets/patelprashant/employee-attrition
#2. Perform data preprocessing
#3. Model building using logistic regression and decision tree
#4. Calculate Performance metrics
```

```
In [3]: #For linear algebra
import numpy as np
#For data processing
import pandas as pd
#For Visualization
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [19]: df=pd.read_csv("Employee-Attrition.csv")
print(df.head)
```

```
<bound method NDFrame.head of
      Department \
0      41      Yes      Travel_Rarely      1102      Sales
1      49      No      Travel_Frequently      279      Research & Development
2      37      Yes      Travel_Rarely      1373      Research & Development
3      33      No      Travel_Frequently      1392      Research & Development
4      27      No      Travel_Rarely      591      Research & Development
...      ...      ...      ...      ...      ...
1465    36      No      Travel_Frequently      884      Research & Development
1466    39      No      Travel_Rarely      613      Research & Development
1467    27      No      Travel_Rarely      155      Research & Development
1468    49      No      Travel_Frequently      1023      Sales
1469    34      No      Travel_Rarely      628      Research & Development

      DistanceFromHome      Education      EducationField      EmployeeCount \
0              1              2      Life Sciences              1
1              8              1      Life Sciences              1
2              2              2              Other              1
3              3              4      Life Sciences              1
4              2              1              Medical              1
...      ...      ...      ...      ...
1465          23              2              Medical              1
1466           6              1              Medical              1
1467           4              3      Life Sciences              1
1468           2              3              Medical              1
1469           8              3              Medical              1

      EmployeeNumber      ...      RelationshipSatisfaction      StandardHours \
0              1      ...              1              80
1              2      ...              4              80
2              4      ...              2              80
3              5      ...              3              80
4              7      ...              4              80
...      ...      ...      ...      ...
1465          2061      ...              3              80
1466          2062      ...              1              80
1467          2064      ...              2              80
1468          2065      ...              4              80
1469          2068      ...              1              80

      StockOptionLevel      TotalWorkingYears      TrainingTimesLastYear \
0              0              8              0
1              1             10              3
2              0              7              3
3              0              8              3
4              1              6              3
```

```

...      ...      ...
1465      1      17      3
1466      1      9      5
1467      1      6      0
1468      0      17      3
1469      0      6      3

```

```

      WorkLifeBalance  YearsAtCompany  YearsInCurrentRole  \
0      1      6      4
1      3      10      7
2      3      0      0
3      3      8      7
4      3      2      2
...      ...      ...
1465      3      5      2
1466      3      7      7
1467      3      6      2
1468      2      9      6
1469      4      4      3

```

```

      YearsSinceLastPromotion  YearsWithCurrManager
0      0      5
1      1      7
2      0      0
3      3      0
4      2      2
...      ...
1465      0      3
1466      1      7
1467      0      3
1468      0      8
1469      1      2

```

[1470 rows x 35 columns]>

In [5]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                  1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                            1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                            1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                        1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                       1470 non-null   int64
17  MaritalStatus                        1470 non-null   object
18  MonthlyIncome                        1470 non-null   int64
19  MonthlyRate                          1470 non-null   int64
20  NumCompaniesWorked                   1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                    1470 non-null   int64

```


24	PerformanceRating	1470	non-null	int64
25	RelationshipSatisfaction	1470	non-null	int64
26	StandardHours	1470	non-null	int64
27	StockOptionLevel	1470	non-null	int64
28	TotalWorkingYears	1470	non-null	int64
29	TrainingTimesLastYear	1470	non-null	int64
30	WorkLifeBalance	1470	non-null	int64
31	YearsAtCompany	1470	non-null	int64
32	YearsInCurrentRole	1470	non-null	int64
33	YearsSinceLastPromotion	1470	non-null	int64
34	YearsWithCurrManager	1470	non-null	int64

dtypes: int64(26), object(9)

memory usage: 402.1+ KB


```
In [6]: #For checking Missing values
plt.figure(figsize=(10,4))
sns.heatmap(df.isnull(), yticklabels=False, cbar=False, cmap='viridis')
```

Out[6]: <Axes: >




```
In [7]: sns.set_style('darkgrid')
sns.countplot(x='Attrition', data=df)
```

Out[7]: <Axes: xlabel='Attrition', ylabel='count'>




```
In [8]: sns.lmplot(x='Age', y='DailyRate', hue='Attrition', data=df)
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x224bb84bfd0>
```



```
In [9]: plt.figure(figsize=(10,6))
sns.boxplot(y='MonthlyIncome', x='Attrition', data=df)
```

Out[9]: <Axes: xlabel='Attrition', ylabel='MonthlyIncome'>



```
In [10]: df.drop(['EmployeeCount', 'StandardHours', 'EmployeeNumber', 'Over18'], axis=1, inplace=True)
print(df.shape)
```

(1470, 31)

```
In [11]: y=df.iloc[:,1]
x=df
x.drop('Attrition', axis=1, inplace=True)
```

```
In [12]: from sklearn.preprocessing import LabelEncoder
lb=LabelEncoder()
y=lb.fit_transform(y)
```

```
In [13]: dum_BusinessTravel = pd.get_dummies(df['BusinessTravel'], prefix='BusinessTravel')
dum_Department = pd.get_dummies(df['Department'], prefix='Department')
dum_EducationField = pd.get_dummies(df['EducationField'], prefix='EducationField') # Fi
dum_Gender = pd.get_dummies(df['Gender'], prefix='Gender', drop_first=True)
dum_JobRole = pd.get_dummies(df['JobRole'], prefix='JobRole')
dum_MaritalStatus = pd.get_dummies(df['MaritalStatus'], prefix='MaritalStatus')
dum_OverTime = pd.get_dummies(df['OverTime'], prefix='OverTime', drop_first=True)

# Adding these dummy variable to input X
X = pd.concat([x, dum_BusinessTravel, dum_Department, dum_EducationField, dum_Gender, du

# Removing the categorical data
X.drop(['BusinessTravel', 'Department', 'EducationField', 'Gender', 'JobRole', 'MaritalS

print(X.shape)
print(y.shape)
```

(1470, 49)

(1470,)

```
In [14]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state
```

```

from sklearn.neighbors import KNeighborsClassifier
neighbors = []
cv_scores = []

```

```

In [15]: from sklearn.model_selection import cross_val_score

# Perform 10 fold cross-validation
for k in range(1, 40, 2):
    neighbors.append(k)
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy')
    cv_scores.append(scores.mean())


# Calculate the error rate
error_rate = [1-x for x in cv_scores]

# Determine the best k
optimal_k = neighbors[error_rate.index(min(error_rate))]
print('The optimal number of neighbors is %d' % optimal_k)

# Plot misclassification error versus k
plt.figure(figsize=(10, 6))
plt.plot(range(1, 40, 2), error_rate, color='blue', linestyle='dashed', marker='o', mark
plt.xlabel('Number of neighbors')
plt.ylabel('Misclassification Error')
plt.show()

```

The optimal number of neighbors is 7



```

In [16]: from sklearn.model_selection import cross_val_predict, cross_val_score
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import confusion_matrix
def print_score(clf, X_train, y_train, X_test, y_test, train = True):
    if train:
        print("Train Result:")
        print("_____")
        print("Classification Report: \n {}".format(classification_report(y_train, clf
        print("Confusion Matrix: \n 0\n".format(confusion_matrix(y_train, clf.predict(X_
        res = cross_val_score(clf, X_train, y_train, cv = 10, scoring = 'accuracy')

```

```

print("Average Accuracy: \t {0:.4f}".format(np.mean(res)))
print("Accuracy SD: \t\t {0:.4f}".format(np.std(res)))
print("accuracy score: {0:.4f}\n".format(accuracy_score(y_train, clf.predict(X_train)))
print("
elif train == False:
    print("Test Result:")
    print("
    print("Classification Report: \n {} \n".format(classification_report(y_test, clf.predict(X_test)))
    print("Confusion Matrix: \n {} \n".format(confusion_matrix(y_test, clf.predict(X_test)))
    print("accuracy score: {0:.4f}\n".format(accuracy_score(y_test, clf.predict(X_test)))
    print("
knn=KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train, y_train)
print_score(knn, X_train, y_train, X_test, y_test, train=True)
print_score(knn, X_train, y_train, X_test, y_test, train=False)

```

Train Result:

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.99	0.92	922
1	0.83	0.19	0.32	180
accuracy			0.86	1102
macro avg	0.85	0.59	0.62	1102
weighted avg	0.86	0.86	0.82	1102

Confusion Matrix:

0

Average Accuracy: 0.8421
Accuracy SD: 0.0148
accuracy score: 0.8621

Test Result:

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.96	0.90	311
1	0.14	0.04	0.06	57
accuracy			0.82	368
macro avg	0.49	0.50	0.48	368
weighted avg	0.74	0.82	0.77	368

Confusion Matrix:

```
[[299 12]
 [ 55  2]]
```

accuracy score: 0.8179

```

In [17]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression

```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

```

```

In [42]: #separate the target variable (Attrition) from features
X=df.drop('Attrition', axis=1)
y=df['Attrition']

#Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

#Define categorical and numerical columns
categorical_columns=['Department', 'JobRole', 'MaritalStatus', 'Gender', 'OverTime']
numerical_columns=['Age', 'DailyRate', 'HourlyRate', 'MonthlyRate', 'NumCompaniesWorked',

#Create transformers for preprocessing
categorical_transformer=Pipeline(steps=[('onehot', OneHotEncoder(handle_unknown='ignore')
numerical_transformer=Pipeline(steps=[('scaler', StandardScaler())])

#combine transformers using a ColumnTransformer
preprocessor=ColumnTransformer(transformers=[('num', numerical_transformer, numerical_co

#Create pipelines for models
logistic_pipeline=Pipeline(steps=[('preprocessor', preprocessor), ('classifier', Logistic
tree_pipeline=Pipeline(steps=[('preprocessor', preprocessor), ('classifier', DecisionTre

```

```

In [43]: from sklearn.preprocessing import OneHotEncoder

#Create a onehot encoder object
encoder=OneHotEncoder(handle_unknown='ignore')

```

```

In [44]: import pandas as pd
from scipy.sparse import csr_matrix

#create a sparse matrix
X=csr_matrix([[1,2], [3,4]])

#convert the sparse matrix to a dense matrix
X_dense=X.todense()

#create a pandas dataframe from the dense matrix
X_df=pd.DataFrame(X_dense)

#print the dataframe
print(X_df)

```

```

    0  1
0  1  2
1  3  4

```

```

In [45]: from scipy.sparse import csr_matrix

```

```

In [40]: #convert the sparse matrices to dense matrices
X_train_dense=X_train.any()
X_test_dense=X_test.any()

#creat pandas DataFrames from the dense matrices
X_train_df=pd.DataFrame(X_train_dense)
X_test_df=pd.DataFrame(X_test_dense)

#Now you can use strings to specify columns
#X_train_encoded=encoder.fit_transform(X_train_df)
#X_test_df=pd.DataFrame(X_test_dense)
X_train_encoded = encoder.transform(X_train)
X_test_df = encoder.transform(X_test)

```



```
#Fit the models
logistic_pipeline.fit(X_train_encoded, y_train)
tree_pipeline.fit(X_train_encoded, y_train)
```

```
-----
NotFittedError                                Traceback (most recent call last)
Cell In[40], line 12
      7 X_test_df=pd.DataFrame(X_test_dense)
      9 #Now you can use strings to specify columns
     10 #X_train_encoded=encoder.fit_transform(X_train_df)
     11 #X_test_df=pd.DataFrame(X_test_dense)
--> 12 X_train_encoded = encoder.transform(X_train)
     13 X_test_df = encoder.transform(X_test)
     15 #Fit the models

File ~\anaconda3\Lib\site-packages\sklearn\utils\_set_output.py:157, in _wrap_method_out
put.<locals>.wrapped(self, X, *args, **kwargs)
     155 @wraps(f)
     156 def wrapped(self, X, *args, **kwargs):
--> 157     data_to_wrap = f(self, X, *args, **kwargs)
     158     if isinstance(data_to_wrap, tuple):
     159         # only wrap the first output for cross decomposition
     160         return_tuple = (
     161             _wrap_data_with_container(method, data_to_wrap[0], X, self),
     162             *data_to_wrap[1:],
     163         )

File ~\anaconda3\Lib\site-packages\sklearn\preprocessing\_encoders.py:1013, in OneHotEnc
oder.transform(self, X)
     994 def transform(self, X):
     995     """
     996     Transform X using one-hot encoding.
     997
     (... )
    1011         returned.
    1012     """
-> 1013     check_is_fitted(self)
    1014     transform_output = _get_output_config("transform", estimator=self)["dense"]
    1015     if transform_output == "pandas" and self.sparse_output:

File ~\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1461, in check_is_fitted
(estimator, attributes, msg, all_or_any)
    1458     raise TypeError("%s is not an estimator instance." % (estimator))
    1460 if not _is_fitted(estimator, attributes, all_or_any):
-> 1461     raise NotFittedError(msg % {"name": type(estimator).__name__})

NotFittedError: This OneHotEncoder instance is not fitted yet. Call 'fit' with appropria
te arguments before using this estimator.
```

```
In [46]: import pandas as pd
from sklearn.preprocessing import LabelEncoder

# Create a LabelEncoder for each categorical column
encoders = {}
for column in X_train.columns:
    if X_train[column].dtype == 'object':
        encoder = LabelEncoder()
        encoder.fit(pd.concat([X_train[column], X_test[column]]))
        encoders[column] = encoder

# Encode categorical variables in both X_train and X_test
X_train_encoded = X_train.copy()
X_test_encoded = X_test.copy()
```

```

for column, encoder in encoders.items():
    if column in X_train_encoded.columns:
        X_train_encoded[column] = encoder.transform(X_train_encoded[column])
    if column in X_test_encoded.columns:
        X_test_encoded[column] = encoder.transform(X_test_encoded[column])

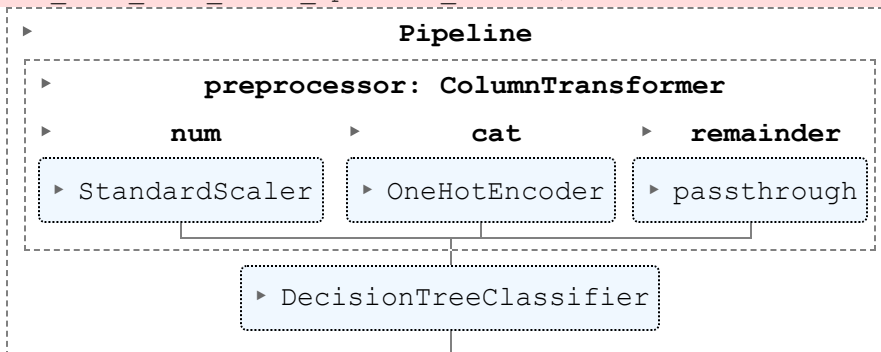
# Fit the models
logistic_pipeline.fit(X_train_encoded, y_train)
tree_pipeline.fit(X_train_encoded, y_train)

```

C:\Users\DELL\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

Out[46]:



In [49]:

```

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

#Encode the categorical features
categorical_features=['Gender', 'Department']
for feature in categorical_features:
    df=pd.get_dummies(df, drop_first=True, columns=[feature])

#split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(df[['Age', 'DailyRate', 'Gender', 'D

```

KeyError

Traceback (most recent call last)

Cell In[49], line 9

```

6     df=pd.get_dummies(df, drop_first=True, columns=[feature])
8 #split the data into train and test sets
----> 9 X_train, X_test, y_train, y_test = train_test_split(df[['Age', 'DailyRate', 'Gen
der', 'Department']], df['Attrition'], test_size=0.25, random_state=42)

```

File ~\anaconda3\Lib\site-packages\pandas\core\frame.py:3813, in DataFrame.__getitem__(self, key)

```

3811     if is_iterator(key):
3812         key = list(key)
-> 3813     indexer = self.columns._get_indexer_strict(key, "columns")[1]
3815 # take() does not accept boolean indexers
3816 if getattr(indexer, "dtype", None) == bool:

```

File ~\anaconda3\Lib\site-packages\pandas\core\indexes\base.py:6070, in Index._get_indexer_strict(self, key, axis_name)

```

6067 else:
6068     keyarr, indexer, new_indexer = self._reindex_non_unique(keyarr)
-> 6070 self._raise_if_missing(keyarr, indexer, axis_name)
6072 keyarr = self.take(indexer)
6073 if isinstance(key, Index):
6074     # GH 42790 - Preserve name from an Index

```

```
File ~\anaconda3\Lib\site-packages\pandas\core\indexes\base.py:6133, in Index._raise_if_missing(self, key, indexer, axis_name)
    6130     raise KeyError(f"None of [{key}] are in the [{axis_name}]")
    6132 not_found = list(ensure_index(key)[missing_mask.nonzero()[0]].unique())
-> 6133 raise KeyError(f"{not_found} not in index")

KeyError: "['Gender', 'Department'] not in index"
```

In []: